

# Communicators and Ranks

Our first MPI for python example will simply import MPI from the mpi4py package, create a *communicator* and get the *rank* of each process:

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

print('My rank is ',rank)
```

Save this to a file call `comm.py` and then run it:

```
mpirun -n 4 python comm.py
```

Here we used the default communicator named `MPI.COMM_WORLD`, which consists of all the processors. For many MPI codes, this is the main communicator that you will need. However, you can create custom communicators using subsets of the processors in `MPI.COMM_WORLD`. See the documentation for more info.

## Point-to-Point Communication

Now we will look at how to pass data from one process to another. Here is a very simple example where we pass a dictionary from process 0 to process 1:

```
from mpi4py import MPI
import numpy

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    data = {'a': 7, 'b': 3.14}
    comm.send(data, dest=1)

elif rank == 1:
    data = comm.recv(source=0)
    print('On process 1, data is ',data)
```

Here we sent a dictionary, but you could also send an integer with a similar code:

```
from mpi4py import MPI
import numpy

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    idata = 1
    comm.send(idata, dest=1)
elif rank == 1:
    idata = comm.recv(source=0)
    print('On process 1, data is ',idata)
```

Note how `comm.send` and `comm.recv` have lower case `s` and `r`.

Now let's look at a more complex example where we send a numpy array:

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    # in real code, this section might
    # read in data parameters from a file
    numData = 10
    comm.send(numData, dest=1)

    data = np.linspace(0.0,3.14,numData)
    comm.Send(data, dest=1)

elif rank == 1:
```

```

numData = comm.recv(source=0)
print('Number of data to receive: ', numData)

data = np.empty(numData, dtype='d') # allocate space to receive the array
comm.Recv(data, source=0)

print('data received: ', data)

```

Note how `comm.Send` and `comm.Recv` used to send and receive the numpy array have upper case `S` and `R`.

## Collective Communication

### Broadcasting:

Broadcasting takes a variable and sends an exact copy of it to all processes on a communicator. Here are some examples:

Broadcasting a dictionary:

```

from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    data = {'key1' : [1,2, 3],
            'key2' : ('abc', 'xyz')}
else:
    data = None

data = comm.bcast(data, root=0)
print('Rank: ', rank, ', data: ', data)

```

You can broadcasting a numpy array using the `Bcast` method (again note the capital `B`). Here we will modify the point-to-point code from above to instead broadcast the array `data` to all processes in the communicator (rather than just from process 0 to 1):

```

from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    # create a data array on process 0
    # in real code, this section might
    # read in data parameters from a file
    numData = 10
    data = np.linspace(0.0,3.14,numData)
else:
    numData = None

# broadcast numData and allocate array on other ranks:
numData = comm.bcast(numData, root=0)
if rank != 0:
    data = np.empty(numData, dtype='d')

comm.Bcast(data, root=0) # broadcast the array from rank 0 to all others

print('Rank: ',rank, ', data received: ',data)

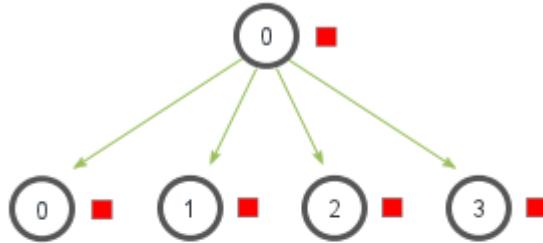
```

## Scattering:

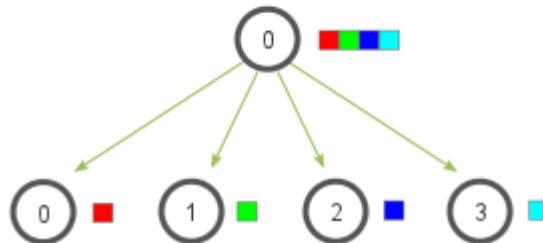
Scatter takes an array and distributes contiguous sections of it across the ranks of a

communicator. Here's an image from <http://mpitutorial.com> that illustrates the difference

MPI\_Bcast



MPI\_Scatter



between a broadcast and scatter:

Let's try an example now.

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
size = comm.Get_size() # new: gives number of ranks in comm
rank = comm.Get_rank()

numDataPerRank = 10
data = None
if rank == 0:
    data = np.linspace(1,size*numDataPerRank,numDataPerRank*size)
    # when size=4 (using -n 4), data = [1.0:40.0]

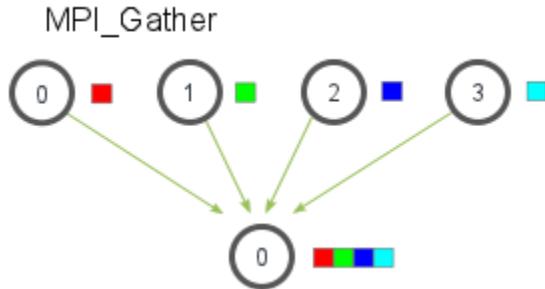
recvbuf = np.empty(numDataPerRank, dtype='d') # allocate space for recvbuf
comm.Scatter(data, recvbuf, root=0)

print('Rank: ',rank, ', recvbuf received: ',recvbuf)
```

In this example, the rank 0 process created the array `data`. Since this is just a toy example, we made `data` be a simple `linspace` array, but in a research code the data might have been read in from a file, or generated by a previous part of the workflow. `data` is then scattered to all the ranks (including rank 0) using `comm.Scatter`. Note that we first had to initialize (or allocate) the receiving buffer array `recvbuf`.

## Gathering:

The reverse of a `scatter` is a `gather`, which takes subsets of an array that are distributed across the ranks, and *gathers* them back into the full array. Here's an image from <http://mpitutorial.com> that illustrates this graphically:



For an example, here is a code that does the reverse of the scatter example above.

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
size = comm.Get_size()
rank = comm.Get_rank()

numDataPerRank = 10
sendbuf = np.linspace(rank*numDataPerRank+1,(rank+1)*numDataPerRank,numDataPerRank)
print('Rank: ',rank, ', sendbuf: ',sendbuf)

recvbuf = None
if rank == 0:
    recvbuf = np.empty(numDataPerRank*size, dtype='d')

comm.Gather(sendbuf, recvbuf, root=0)
```

```
if rank == 0:  
    print('Rank: ',rank, ', recvbuf received: ',recvbuf)
```

## Reduction:

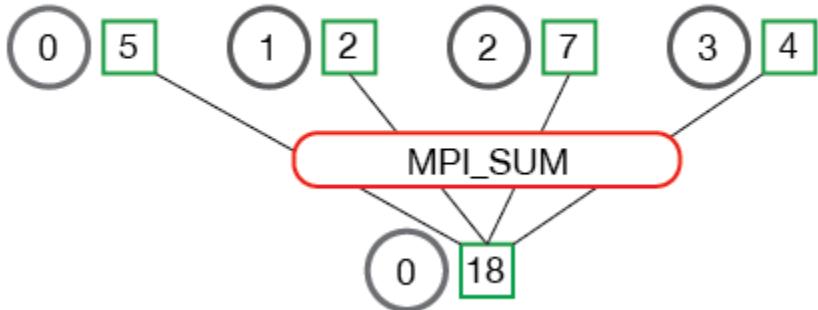
The MP `reduce` operation takes values in from an array on each process and reduces them to a single result on the root process. This is essentially like having a somewhat complicated send command from each process to the root process, and then having the root process perform the reduction operation. Thankfully, MPI `reduce` does all this with one concise command.

For numpy arrays, the syntax is

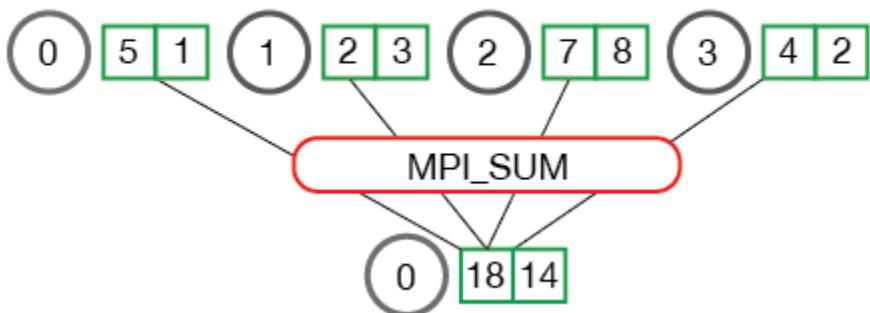
~~~python comm.Reduce(send\_data, recv\_data, op=, root=0) ~~~ where `send_data` is the data being sent from all the processes on the communicator and `recv_data` is the array on the root process that will receive all the data. Note that the send and recv arrays have the same size. All the copies of data in `send_data` will be reduced according to the specified `<operation>`. A few commonly used operations are: - MPI\_SUM - Sums the elements. - MPI\_PROD - Multiplies all elements. - MPI\_MAX - Returns the maximum element. - MPI\_MIN - Returns the minimum element.

Here's a few more useful images from <http://mpitutorial.com> that illustrate the reduce step graphically:

MPI\_Reduce



MPI\_Reduce



And here is a code example:

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

# Create some np arrays on each process:
# For this demo, the arrays have only one
# entry that is assigned to be the rank of the processor
value = np.array(rank, 'd')

print(' Rank: ',rank, ' value = ', value)

# initialize the np arrays that will store the results:
value_sum      = np.array(0.0, 'd')
value_max      = np.array(0.0, 'd')

# perform the reductions:
comm.Reduce(value, value_sum, op=MPI.SUM, root=0)
comm.Reduce(value, value_max, op=MPI.MAX, root=0)

if rank == 0:
    print(' Rank 0: value_sum =     ',value_sum)
    print(' Rank 0: value_max =     ',value_max)
```