



Parallel & Distributed Computing

Lecture week 11:

Distributed Memory Programming

Farhad M. Riaz

Farhad.Muhammad@numl.edu.pk

Department of Computer Science
NUML, Islamabad

Communicators and Ranks

```
from mpi4py import MPI
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
print('My rank is ',rank)
```

Point-to-Point Communication

```
from mpi4py import MPI
import numpy
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
if rank == 0:
    data = {'a': 7, 'b': 3.14}
    comm.send(data, dest=1)
elif
    rank == 1: data = comm.recv(source=0)
print('On process 1, data is ',data)
```

Collective Communication Broadcasting:

- Broadcasting takes a variable and sends an exact copy of it to all processes on a communicator.

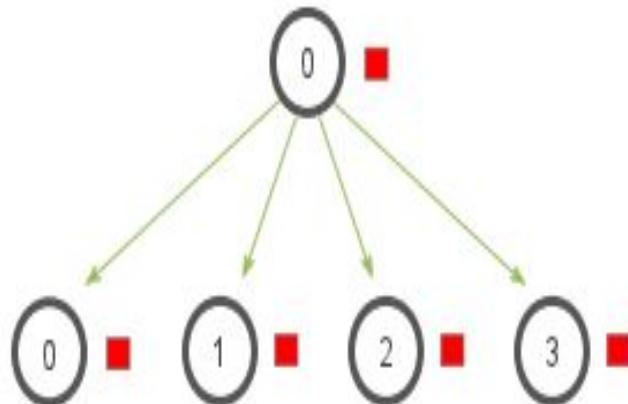
```
from mpi4py import MPI
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
if rank == 0:
    data = {'key1': [1,2, 3], 'key2': ( 'abc', 'xyz')}
else:
    data = None
    data = comm.bcast(data, root=0)
print('Rank: ',rank,' ,data: ' ,data)
```

```
from mpi4py import MPI
import numpy as np
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
if rank == 0:
    # create a data array on process 0
    # in real code, this section might
    # read in data parameters from a file
    numData = 10
    data = np.linspace(0.0,3.14,numData)
else:
    numData = None
    # broadcast numData and allocate array on other ranks:
    numData = comm.bcast(numData, root=0)
if rank != 0:
    data = np.empty(numData, dtype='d')
    comm.Bcast(data, root=0)
    # broadcast the array from rank 0 to all others
print('Rank: ',rank, ', data received: ',data)
```

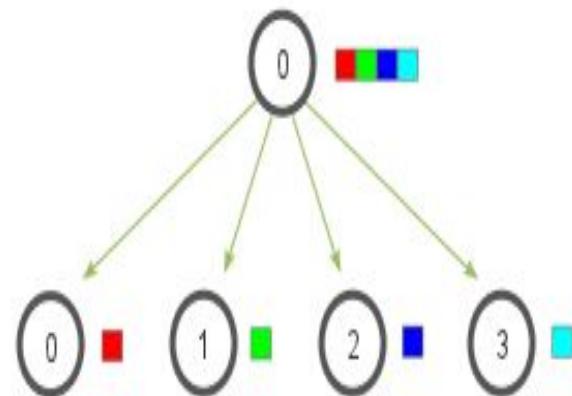
Scattering:

- Scatter takes an array and distributes contiguous sections of it across the ranks of a communicator

MPI_Bcast



MPI_Scatter

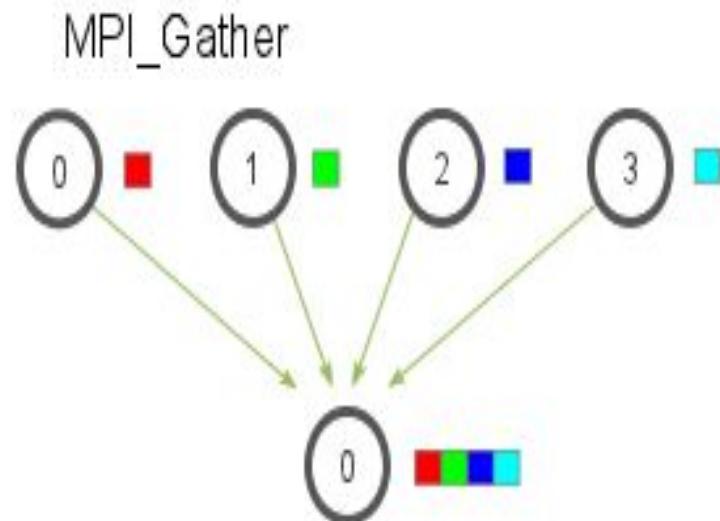


Example

```
from mpi4py import MPI
import numpy as np
comm = MPI.COMM_WORLD
size = comm.Get_size()
    # new: gives number of ranks in comm
rank = comm.Get_rank()
numDataPerRank = 10
data = None
if rank == 0:
    data=np.linspace(1,size*numDataPerRank,numDataPerRank*size)
        # when size=4 (using -n 4), data = [1.0:40.0]
recvbuf = np.empty(numDataPerRank, dtype='d')
    # allocate space for recvbuf
comm.Scatter(data, recvbuf, root=0)
print('Rank: ',rank, ', recvbuf received: ',recvbuf)
```

Gathering:

The reverse of a **scatter** is a **gather**, which takes subsets of an array that are distributed across the ranks, and *gathers* them back into the full array



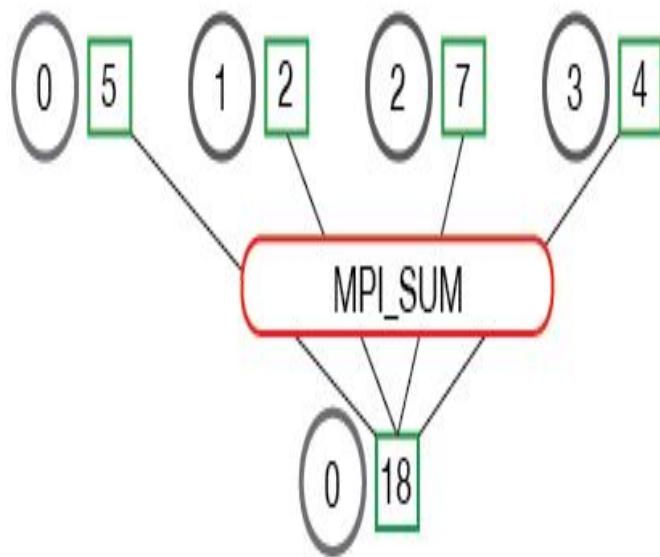
Example

```
from mpi4py import MPI
import numpy as np
comm = MPI.COMM_WORLD
size = comm.Get_size()
rank = comm.Get_rank()
numDataPerRank = 10
Sendbuf=np.linspace(rank*numDataPerRank+1,(rank+1)*numDataPerRank,numDataPerRank)
print('Rank: ',rank, ', sendbuf: ',sendbuf)
recvbuf = None
if rank == 0:
    recvbuf = np.empty(numDataPerRank*size, dtype='d') comm.Gather(sendbuf, recvbuf,
root=0)
if rank == 0:
    print('Rank: ',rank, ', recvbuf received: ',recvbuf)
```

Reduction:

- The MP **reduce** operation takes values in from an array on each process and reduces them to a single result on the root process.
- This is essentially like having a somewhat complicated send command from each process to the root process, and then having the root process perform the reduction operation.

MPI_Reduce



MPI_Reduce

